

## See chapter 4 in Regression and Other Stories.

Widen the notebook.

```
•  
• html"""  
• <style>  
•     main {  
•         margin: 0 auto;  
•         max-width: 2000px;  
•         padding-left: max(160px, 10%);  
•         padding-right: max(160px, 10%);  
•     }  
• </style>  
• """
```

```
• using Pkg ✓ , DrWatson ✓
```

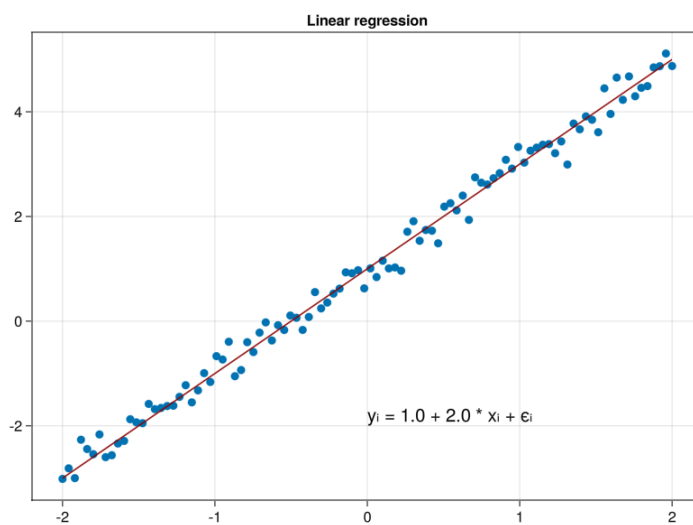
A typical set of Julia packages to include in notebooks.

```
• begin  
•     # Specific to this notebook  
•     using GLM ✓  
•  
•     # Specific to ROSTuringPluto  
•     using Optim ✓  
•     using Logging ✓  
•     using Turing ✓  
•  
•     # Graphics related  
•     using GLMakie ✓  
•  
•     # Common data files and functions  
•     using RegressionAndOtherStories ✓  
•     import RegressionAndOtherStories: link  
•  
•     Logging.disable_logging(Logging.Warn)  
• end;
```

Replacing docs for `RegressionAndOtherStories.tr  
DataFrame, AbstractString}` in module `Regressio

## 4.1 Sampling distributions and generative models.

```
• begin
•   Random.seed!(1)
•   a = 1.0
•   b = 2.0
•   x = LinRange(-2, 2, 100)
•   y = a .+ b .* x .+ rand(Normal(0.0,
•   0.2), 100)
• end;
```



```
• let
•   f = Figure()
•   ax = Axis(f[1, 1]; title="Linear
•   regression")
•   scatter!(x, y)
•   lines!(x, a .+ b .* x; color=:darkred)
•   annotations!("yi = 1.0 + 2.0 * xi +
•   εi", position=(0, -2), fontsize=20)
•   current_figure()
• end
```

ppl4\_1 (generic function with 2 methods)

```
• @model function ppl4_1(x, y)
•   a ~ Normal(0, 1.5)
•   b ~ Normal(1.0, 1.5)
•   σ ~ Exponential(1)
•   μ = a .+ b .* x
•   for i in eachindex(y)
•       y[i] ~ Normal(μ[i], σ)
•   end
• end
```

	parameters	mean	std	naive_
1	:a	1.00165	0.0217711	0.000342
2	:b	1.97541	0.0187714	0.000296
3	: $\sigma$	0.218379	0.0158395	0.000256

```

• begin
•   m4_1t = ppl4_1(x, y)
•   chns4_1t = sample(m4_1t, NUTS(),
•   MCMCThreads(), 1000, 4)
•   describe(chns4_1t)
• end


```

	parameters	median	mad_sd	mean	st
1	"a"	1.002	0.022	1.002	0.02
2	"b"	1.976	0.019	1.975	0.01
3	" $\sigma$ "	0.218	0.016	0.218	0.01

```

• begin
•   post4_1t = DataFrame(chns4_1t)[: , 3:5]
•   ms4_1t = model_summary(post4_1t,
•   names(post4_1t))
• end


```



```

• let
•   f = Figure()
•   ax = Axis(f[1, 1]; title="Sampling
•   distribution of a")
•   density!(post4_1t.a)
•   ax = Axis(f[1, 2]; title="Sampling
•   distribution of b")
•   density!(post4_1t.b)
•   ax = Axis(f[1, 3]; title="Sampling
•   distribution of sigma")
•   density!(post4_1t.σ)
•   current_figure()
end

```




```

• let
•   f = Figure()
•   ax = Axis(f[1, 1]; title="Linear
•   regression")
•   scatter!(x, y)
•   lines!(x, ms4_1t["a", "median"] .+
•   ms4_1t["b", "median"] .* x;
•   color=:darkred)
•   mean_a = round(ms4_1t["a", "mean"];
•   digits=2)
•   mean_b = round(ms4_1t["b", "mean"];
•   digits=2)
•   mean_σ = round(ms4_1t["σ", "mean"];
•   digits=2)
•   annotations!("y = $(mean_a) + $(mean_b)
•   * x + $(mean_σ)", position=(0, -2),
•   fontsize=20)
•   current_figure()
• end

```

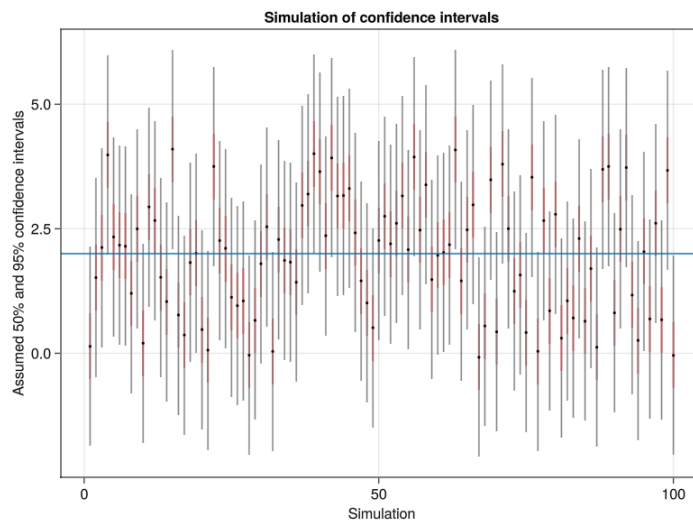
## 4.2 Estimates, standard errors, and confidence intervals.



```

• let
•   f = Figure()
•   ax = Axis(f[1, 1]; title="Sampling
•   distribution of b (revisited)")
•   b̂ = ms4_1t["b", "median"]
•   σ̂ = ms4_1t["b", "std"]
•   x = LinRange(b̂ - 4σ̂, b̂ + 4σ̂, 100)
•   y = pdf.(Normal(b̂, σ̂), x)
•   ylims!(ax, [0, maximum(y) + 1.0])
•   ax.xticks = b̂ - 3σ̂ : σ̂ : b̂ + 3σ̂
•   ax.xtickformat = xs -> ["$(i) s.e." for
•   i in -3:3]
•   lines!(x, y)
•   vlines!(ax, b̂;
•   ymax=maximum(y)/(maximum(y) + 1.0),
•   color=:grey)
•   vlines!(ax, b̂-σ̂; ymax=pdf.(Normal(b̂,
•   σ̂), b̂-σ̂)/(maximum(y) + 1.0),
•   color=:grey)
•   vlines!(ax, b̂+σ̂; ymax=pdf.(Normal(b̂,
•   σ̂), b̂+σ̂)/(maximum(y) + 1.0),
•   color=:grey)
•   annotations!("b ± 1 s.e.", position=
•   (b̂-0.008, 7.5), textsize=20)
•   x1 = range(b̂ - σ̂, b̂ + σ̂; length=60)
•   band!(x1, fill(0, length(x1)), pdf.
•   (Normal(b̂, σ̂), x1); color = (:blue,
•   0.25))
•   f
• end

```



```

let
  n = 100
  b = 2.0

  f = Figure()
  ax = Axis(f[1,1]; title="Simulation of
confidence intervals",
xlabel="Simulation",
ylabel="Assumed 50% and 95%
confidence intervals")

  x = 1:n
  y = [rand(Uniform(b - 2.1 , b + 2.1), 1)
[1] for i in 1:n]

  # Assumed s.e. = 1.0
  lowererrors = fill(0.66, n)
  highererrors = fill(2, n)

  errorbars!(x, y, lowererrors, color =
:red) # same low and high error
  errorbars!(x, y, highererrors, color =
:grey) # same low and high error

  scatter!(x, y, markersize = 3, color =
:black)
  hlines!(ax, [2])

  f
end

```

```
► (estimate = 0.7, se = 0.0144914, int_95 = [0.67
```

```
• let  
•   n = 1000  
•   yes = 700  
•   no = n - yes  
•   est = yes/n  
•   se = sqrt(est * (1 - est)/n)  
•  
•   (estimate = est, se = se, int_95 = est  
•     .+ quantile.(Normal(0, 1), [0.025,  
•     0.975]) * se)  
end
```

```
► (estimate = 35.8, se = 0.734847, int_50 = [35.2
```


```
• let  
•   y = [35, 34, 38, 35, 37]  
•   n = length(y)  
•   est = mean(y)  
•   se = std(y)/sqrt(n)  
•   int_50 = est .+ quantile.(TDist(n-1),  
•   [0.25, 0.75]) * se  
•   int_95 = est .+ quantile.(TDist(n-1),  
•   [0.025, 0.975]) * se  
•  
•   (estimate = est, se = se, int_50 =  
•   int_50, int_95 = int_95)  
end
```



```
df_poll =
```

	poll1	poll2	poll3	poll4	poll5
<b>1</b>	2002	10.0	70.0	25.0	5.0
<b>2</b>	2002	5.0	72.0	25.0	3.0
<b>3</b>	2001	10.0	68.0	26.0	6.0
<b>4</b>	2001	5.0	65.0	27.0	8.0
<b>5</b>	2001	2.0	67.0	25.0	8.0
<b>6</b>	2000	8.0	67.0	28.0	5.0
<b>7</b>	2000	6.0	66.0	26.0	8.0
<b>8</b>	2000	2.0	66.0	28.0	6.0
<b>9</b>	1999	5.0	71.0	22.0	7.0
<b>10</b>	1995	9.0	77.0	13.0	10.0
⋮ more					
<b>32</b>	1937	12.0	60.0	33.0	7.0

```
• df_poll = CSV.read(ros_datadir("Death",  
  "polls.csv"), DataFrame)
```



```

let
  f = Figure()
  ax = Axis(f[1, 1]; title="Death penalty
  opinions", xlabel="Year",
  ylabel="Percentage support for the
  death penalty")
  scatter!(df_poll.year, df_poll.support
  .* 100)
  global err_lims =
  [100(sqrt(df_poll.support[i]*(1-
  df_poll.support[i])/1000)) for i in
  1:nrow(df_poll)]
  errorbars!(df_poll.year, df_poll.support
  .* 100, err_lims, color = :red)
  f
end

```

► [1.3925, 1.38313, 1.41453, 1.43996, 1.40676, 1.4

• err\_lims

## 4.3 Bias and unmodeled uncertainty.

## 4.4 Statistical significance, hypothesis testing, and statistical errors.

**4.5 Problems with the concept of statistical significance.**

**4.6 Example of hypothesis testing: 55,000 residents need your help!**

**4.7 Moving beyond hypothesis testing.**