

# See chapter 1 in Regression and Other Stories.

=====

Widen the cells.

```
• html"""  
• <style>  
•     main {  
•         margin: 0 auto;  
•         max-width: 2000px;  
•         padding-left: max(160px, 10%);  
•         padding-right: max(160px, 10%);  
•     }  
• </style>  
• """
```

A typical set of Julia packages to include in  
notebooks.

```
• using Pkg ✓
```

```
• begin  
•     # Specific to this notebook  
•     using GLM ✓  
•   
•     # Specific to ROSTuringPluto  
•     using Optim ✓  
•     using Logging ✓  
•     using Turing ✓  
•   
•     # Graphics related  
•     using CairoMakie ✓  
•     using AlgebraOfGraphics ✓  
•   
•     # Common data files and functions  
•     using RegressionAndOtherStories ✓  
•     import RegressionAndOtherStories: link  
•   
•     Logging.disable_logging(Logging.Warn)  
• end;
```

## 1.1 The three challenges of statistics.

### Note

It is not common for me to copy from the book but this particular section deserves an exception!

The three challenges of statistical inference are:

1. Generalizing from sample to population, a problem that is associated with survey sampling but actually arises in nearly every application of statistical inference;
2. Generalizing from treatment to control group, a problem that is associated with causal inference, which is implicitly or explicitly part of the interpretation of most regressions we have seen; and
3. Generalizing from observed measurements to the underlying constructs of interest, as most of the time our data do not record exactly what we would ideally like to study.

All three of these challenges can be framed as problems of prediction (for new people or new items that are not in the sample, future outcomes under different potentially assigned treatments, and underlying constructs of interest, if they could be measured exactly).

## 1.2 Why learn regression?

```
hibbs =
```

	year	growth	vote	inc_party_candidate
1	1952	2.4	44.6	"Stevenson"
2	1956	2.89	57.76	"Eisenhower"
3	1960	0.85	49.91	"Nixon"
4	1964	4.21	61.34	"Johnson"
5	1968	3.02	49.6	"Humphrey"
6	1972	3.62	61.79	"Nixon"
7	1976	1.08	48.95	"Ford"
8	1980	-0.39	44.7	"Carter"
9	1984	3.86	59.17	"Reagan"
10	1988	2.27	53.94	"Bush, Sr."
⋮ more				
16	2012	0.95	52.0	"Obama"

```
• hibbs =  
  CSV.read(ros_datadir("ElectionsEconomy",  
    "hibbs.csv"), DataFrame)
```

```
hibbs_lm =  
StatsModels.TableRegressionModel{LinearModel{GLM},  
  vote ~ 1 + growth
```

Coefficients:

	Coef.	Std. Error	t	Pr(> t )
(Intercept)	46.2476	1.62193	28.51	<1e-30
growth	3.06053	0.696274	4.40	0.0004

```
• hibbs_lm = lm(@formula(vote ~ growth),  
  hibbs)
```

```
► [-8.99292, 2.66743, 1.0609, 2.20753, -5.89044, 4.27444]
```

```
• residuals(hibbs_lm)
```

```
2.2744434224582912
```


```
• mad(residuals(hibbs_lm))
```

```
3.635681268522063
```

```
• std(residuals(hibbs_lm))
```

```
► [46.2476, 3.06053]
```

```
• coef(hibbs_lm)
```



```

let
  fig = Figure()
  hibbs.label = string.(hibbs.year)
  xlabel = "Average growth personal
income [%]"
  ylabel = "Incumbent's party vote share"
  let
    title = "Forecasting the election
from the economy"
    ax = Axis(fig[1, 1]; title, xlabel,
ylabel)
    for (ind, yr) in
      enumerate(hibbs.year)
        annotations!("$ (yr)"; position=
(hibbs.growth[ind],
hibbs.vote[ind]), fontsize=10)
    end
  end
  let
    x = LinRange(-1, 4, 100)
    title = "Data and linear fit"
    ax = Axis(fig[1, 2]; title, xlabel,
ylabel)
    scatter!(hibbs.growth, hibbs.vote)
    lines!(x, coef(hibbs_lm)[1] .+
coef(hibbs_lm)[2] .* x;
color=:darkred)
    annotations!("vote = 46.2 + 3.0 *
growth"; position=(0, 41))
  end
  fig
end

```

ppl7\_1 (generic function with 2 methods)

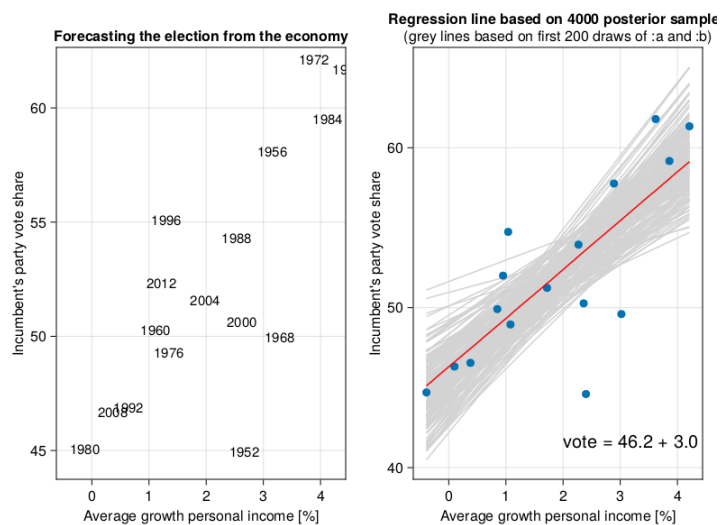
```
• @model function ppl7_1(growth, vote)
•   a ~ Normal(50, 20)
•   b ~ Normal(2, 10)
•   σ ~ Exponential(1)
•   μ = a .+ b .* growth
•   for i in eachindex(vote)
•     vote[i] ~ Normal(μ[i], σ)
•   end
• end
```

```
► [ parameters mean std naive_se
1  :a          46.2939 1.57086 0.0248374
2  :b           3.04094 0.669116 0.0105796
3  :σ           3.61037 0.650903 0.0102917
```

```
• begin
•   m7_1t = ppl7_1(hibbs.growth, hibbs.vote)
•   chns7_1t = sample(m7_1t, NUTS(),
•   MCMCThreads(), 1000, 4)
•   describe(chns7_1t)
• end
```

```
parameters median mad_sd mean st
1 "a"          46.302 1.474 46.294 1.57
2 "b"           3.043 0.628 3.041 0.66
3 "σ"           3.505 0.595 3.61 0.65
```

```
• begin
•   post7_1t = DataFrame(chns7_1t)[: , 3:5]
•   ms7_1t = model_summary(post7_1t,
•   names(post7_1t))
• end
```



```

• let
•   growth_range =
•     LinRange(minimum(hibbs.growth),
•       maximum(hibbs.growth), 200)
•   votes = median.(link(post7_1t, (r,x) ->
•     r.a + x * r.b, growth_range))
•
•   hibbs.label = string.(hibbs.year)
•   xlabel = "Average growth personal
•     income [%]"
•   ylabel="Incumbent's party vote share"
•
•   fig = Figure()
•   let
•     title = "Forecasting the election
• from the economy"
•     plt = data(hibbs) *
•       mapping(:label => verbatim,
•         (:growth, :vote) => Point) *
•       visual(Annotations, fontsize=15)
•     axis = (; title, xlabel, ylabel)
•     draw!(fig[1, 1], plt; axis)
•   end
•
•   ax = Axis(fig[1, 2]; title="Regression
• line based on 4000 posterior samples",
•     subtitle = "(grey lines based on
• first 200 draws of :a and :b)",
•     xlabel, ylabel)
•   for i in 1:200
•     lines!(growth_range, post7_1t.a[i]
•       .+ post7_1t.b[i] .* growth_range,
•       color = :lightgrey)
•   end
•   scatter!(hibbs.growth, hibbs.vote)

```

```
lines!(growth_range, votes, color =
:red)
annotations!("vote = 46.2 + 3.0 *
growth"; position=(2, 41))
fig
end
```

## 1.3 Some examples of regression.

Electric company



	post_test	pre_test	grade	treatment
<b>1</b>	48.9	13.8	1	1
<b>2</b>	70.5	16.5	1	1
<b>3</b>	89.7	18.5	1	1
<b>4</b>	44.2	8.8	1	1
<b>5</b>	77.5	15.3	1	1
<b>6</b>	84.7	15.0	1	1
<b>7</b>	78.9	19.4	1	1
<b>8</b>	86.8	15.0	1	1
<b>9</b>	60.8	11.8	1	1
<b>10</b>	75.7	16.4	1	1
⋮ more				
<b>192</b>	110.0	102.6	4	0

```

• begin
•   electric =
•   CSV.read(ros_datadir("ElectricCompany",
•   "electric.csv"), DataFrame)
•   electric = electric[:, [:post_test,
•   :pre_test, :grade, :treatment]]
•   electric.grade =
•   categorical(electric.grade)
•   electric.treatment =
•   categorical(electric.treatment)
•   electric
• end

```

**A quick look at the overall values of pre\_test and post\_test.**


	variable	mean	min	median	max
1	:post_test	97.1495	44.2	102.3	122.0
2	:pre_test	72.2245	8.8	80.75	119.8
3	:grade	nothing	1	nothing	4
4	:treatment	nothing	0	nothing	1

```
• describe(electric)
```


true

```
• all(completecases(electric)) == true
```

**Post-test density for each grade  
conditioned on treatment.**



```
• let
•   f = Figure()
•   axis = (; width = 150, height = 150)
•   el = data(electric) *
•   mapping(:post_test, col=:grade,
•   color=:treatment)
•   plt = el *
•   AlgebraOfGraphics.histogram(;bins=20) *
•   mapping(row=:treatment)
•   draw!(f[1, 1], plt; axis)
•   f
end
```



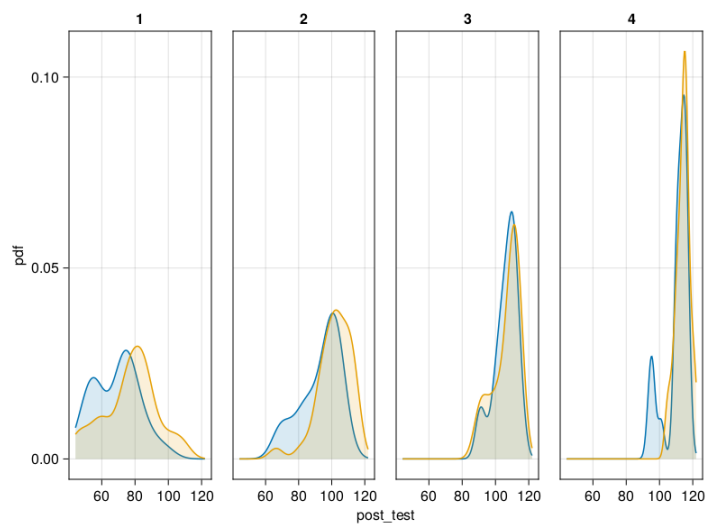
```

let
  f = Figure()
  axis = (; width = 150, height = 150)
  el = data(electric) *
    mapping(:post_test, col=:grade,
            color=:treatment)
  plt = el * AlgebraOfGraphics.density()
    * mapping(row=:treatment)
  draw!(f[1, 1], plt; axis)
f
end

```

### Note


In above cell, as `density()` is exported by both GLMakie and AlgebraOfGraphics, it needs to be qualified.



```

• let
•   f = Figure()
•   el = data(electric) *
•   mapping(:post_test, col=:grade)
•   plt = el * AlgebraOfGraphics.density()
•   * mapping(color=:treatment)
•   draw!(f[1, 1], plt)
•   f
• end

```



```

• let
•   plt = data(electric) * visual(Violin) *
•   mapping(:grade, :post_test,
•   dodge=:treatment, color=:treatment)
•   draw(plt)
• end

```

## Peacekeeping

peace =

	war	cfdate	faildate
1	"Afghanistan-Mujahideen"	8150	8257
2	"Afghanistan-Taliban"	8466	8505
3	"Algeria-FIS/AIS"	10149	12783
4	"Angola"	7820	8319
5	"Angola"	9089	10564
6	"Azerbaijan-N.K."	8643	8678
7	"Azerbaijan-N.K."	8901	12783
8	"Bangladesh-CHT"	8248	12783
9	"Myanmar-Karen"	8153	9282
10	"Myanmar-Karen"	9296	9907
: more			
96	"Yugoslavia-Kosovo"	10751	12783

```
• peace =  
  CSV.read(ros_datadir("PeaceKeeping",  
    "peacekeeping.csv"), missingstring="NA",  
    DataFrame)
```

	variable	mean	min
1	:war	nothing	"Afghanistan-Mujah:
2	:cfdate	8925.1	6985
3	:faildate	10795.8	7074
4	:peacekeepers	0.354167	0
5	:badness	-8.15228	-12.26
6	:delay	5.12177	0.04
7	:censored	0.416667	0

```
• describe(peace)
```

## A quick look at this Dates stuff!

8150

- `peace.cfdate[1]`

1992-04-25T00:00:00

- `DateTime(1992, 4, 25)`

107 days

- `Date(1992, 8, 10) - Date(1992, 4, 25)`

1970-01-01

- `Date(1970,1,1)`

1992-04-25

- `Date(1970,1,1) + Dates.Day(8150)`

8150 days

- `Date(1992, 4, 25) - Date(1970, 1, 1)`

107

- `peace.faildate[1] - peace.cfdate[1]`

```
• begin
•   pks_df = peace[peace.peacekeepers .==
•   1, [:cfdate, :faildate]]
•   nopks_df = peace[peace.peacekeepers .==
•   0, [:cfdate, :faildate]]
end;
```

0.4166666666666667

- `mean(peace.censored)`

64

- `length(unique(peace.war))`

0.5588235294117647

- `mean(peace[peace.peacekeepers .== 1, :censored])`

0.3387096774193548

- `mean(peace[peace.peacekeepers .== 0, :censored])`

1.382

- `mean(peace[peace.peacekeepers .== 1 .&& peace.censored .== 0, :delay])`

1.5153658536585364


- `mean(peace[peace.peacekeepers .== 0 .&& peace.censored .== 0, :delay])`

1.05

- `median(peace[peace.peacekeepers .== 1 .&& peace.censored .== 0, :delay])`

0.59

- `median(peace[peace.peacekeepers .== 0 .&& peace.censored .== 0, :delay])`



```

• let
•   f = Figure()
•   pks = peace[peace.peacekeepers .== 1
•   .&& peace.censored .== 0, :]
•   nopks = peace[peace.peacekeepers .== 0
•   .&& peace.censored .== 0, :]
•
•   for i in 1:2
•       title = i == 1 ? "Peacekeepers" :
•       "No peacekeepers"
•
•       ax = Axis(f[i, 1]; title,
•       xlabel="Years until return to war",
•       ylabel = "Frequency",
•       yminorticksvisible = true,
•       yminorgridvisible = true,
•       yminorticks = IntervalsBetween(8))
•
•       xlims!(ax, [0, 8])
•       hist!(i == 1 ? pks.delay :
•       nopks.delay; bins=20)
•   end
• f
• end

```

### Note


Censored means conflict had not returned until end of observation period (2004).



```

• begin
•   # Filter out missing badness rows.
•   pb = peace[peace.badness .!= missing,
•   :];
•
•   # Delays until return to war for
•   uncensored, peacekeeper cases
•   pks_uc = pb[pb.peacekeepers .== 1 .&&
•   pb.censored .== 0, :delay]
•   # Delays until return to war for
•   censored, peacekeeper cases
•   pks_c = pb[pb.peacekeepers .== 1 .&&
•   pb.censored .== 1, :delay]
•
•   # No peacekeepr cases.
•   nopks_uc = pb[pb.peacekeepers .== 0 .&&
•   pb.censored .== 0, :delay]
•   nopks_c = pb[pb.peacekeepers .== 0 .&&
•   pb.censored .== 1, :delay]
•
•   # Crude measure (:badness) used for
•   assessing situation
•   badness_pks_uc = pb[pb.peacekeepers .==
•   1 .&& pb.censored .== 0,
•   :badness]
•   badness_pks_c = pb[pb.peacekeepers .== 1
•   .&& pb.censored .== 1,
•   :badness]
•   badness_nopks_uc = pb[pb.peacekeepers
•   .== 0 .&& pb.censored .== 0,
•   :badness]
•   badness_nopks_c = pb[pb.peacekeepers
•   .== 0 .&& pb.censored .== 1,
•   :badness]
• end;

```



```

• begin
•     local f = Figure()
•     ax = Axis(f[1, 1], title = "With UN
•         peacekeepers",
•         xlabel = "Pre-treatment measure of
•             problems in country",
•         ylabel = "Delay [yrs] before return
•             to conflict")
•     sca1 = scatter!(badness_pks_uc, pks_uc)
•     sca2 = scatter!(badness_pks_c, pks_c)
•     xlims!(ax, [-13, -2.5])
•     Legend(f[1, 2], [sca1, sca2],
•         ["Uncensored", "Censored"])
•     ax.xticks = ([-12, -4], ["no so bad",
•         "really bad"])
•
•
•     ax = Axis(f[2, 1], title = "Without UN
•         peacekeepers",
•         xlabel = "Pre-treatment measure of
•             problems in country",
•         ylabel = "Delay [yrs] before return
•             to conflict")
•     sca1 = scatter!(badness_nopks_uc,
•         nopks_uc)
•     sca2 = scatter!(badness_nopks_c,
•         nopks_c)
•     xlims!(ax, [-13, -2.5])
•     Legend(f[2, 2], [sca1, sca2],
•         ["Uncensored", "Censored"])
•     ax.xticks = ([-12, -4], ["no so bad",
•         "really bad"])
•
•     f
• end

```

## 1.4 Challenges in building, understanding, and interpreting regression.

### Simple causal

ppl1\_2a (generic function with 2 methods)

```
• @model function ppl1_2a(x, y)
•   a ~ Normal(10, 10)
•   b ~ Normal(10, 10)
•   σ ~ Exponential(1)
•   μ = a .+ b .* x
•   for i in eachindex(x)
•     y[i] ~ Normal(μ[i], σ)
•   end
• end
```

ppl1\_2b (generic function with 2 methods)

```
• @model function ppl1_2b(x_binary, y)
•   a ~ Normal(10, 10)
•   b ~ Normal(10, 10)
•   σ ~ Exponential(1)
•   μ = a .+ b .* x_binary
•   for i in eachindex(x_binary)
•     y[i] ~ Normal(μ[i], σ)
•   end
• end
```

#### Note

Aki Vehtari did not include a seed number in his code.

► [24.3056, 18.7671, 24.9769, 19.2313, 16.5586, 11

```
• begin
•   Random.seed!(123)
•   n = 50
•   x = rand(Uniform(1, 5), n)
•   x_binary = [x[i] < 3 ? 0 : 1 for i in
•     1:n]
•   y = [rand(Normal(10 + 3x[i], 3), 1)[1]
•     for i in 1:n]
• end
```

	iteration	chain	a	b	$\sigma$
1	501	1	9.2363	3.43394	3.21555
2	502	1	9.10217	3.21645	3.0705
3	503	1	8.6518	3.7155	3.34631
4	504	1	8.53133	3.67576	3.39775
5	505	1	9.60928	3.17707	3.53525
6	506	1	8.97448	3.4119	3.25893
7	507	1	8.62075	3.5889	3.50358
8	508	1	7.38018	3.73752	3.3409
9	509	1	7.77257	3.86555	3.3562
10	510	1	9.12777	3.23734	3.81892
	: more				

```

• begin
•   m1_2at = ppl1_2a(x, y)
•   chns1_2at = sample(m1_2at, NUTS(),
•   MCMCThreads(), 1000, 4)
• end

```

	parameters	mean	std	naive_se
1	:a	9.37512	1.36577	0.0215948
2	:b	3.25449	0.4219	0.00667083
3	: $\sigma$	3.46747	0.344304	0.00544392

```

• describe(chns1_2at)

```

	parameters	median	mad_sd	mean	std
1	"a"	9.384	1.355	9.375	1.36
2	"b"	3.256	0.413	3.254	0.42
3	"σ"	3.448	0.34	3.467	0.34

```

• begin
•   post1_2at = DataFrame(chns1_2at)[: , 3:5]
•   ms1_2at = model_summary(post1_2at,
•   names(post1_2at))
• end

```

	iteration	chain	a	b	σ
1	501	1	16.9696	6.72975	3.53264
2	502	1	16.2386	7.87303	3.42438
3	503	1	15.851	6.53625	3.77021
4	504	1	15.9778	6.89418	4.07833
5	505	1	16.6353	6.5867	3.44341
6	506	1	16.3839	6.62497	3.54545
7	507	1	16.4743	6.28868	3.58246
8	508	1	16.4743	6.28868	3.58246
9	509	1	16.9043	5.90824	4.05398
10	510	1	17.076	5.74135	3.86013
	⋮ more				

```

• begin
•   m1_2bt = ppl1_2b(x_binary, y)
•   chns1_2bt = sample(m1_2bt, NUTS(),
•   MCMCThreads(), 1000, 4)
• end

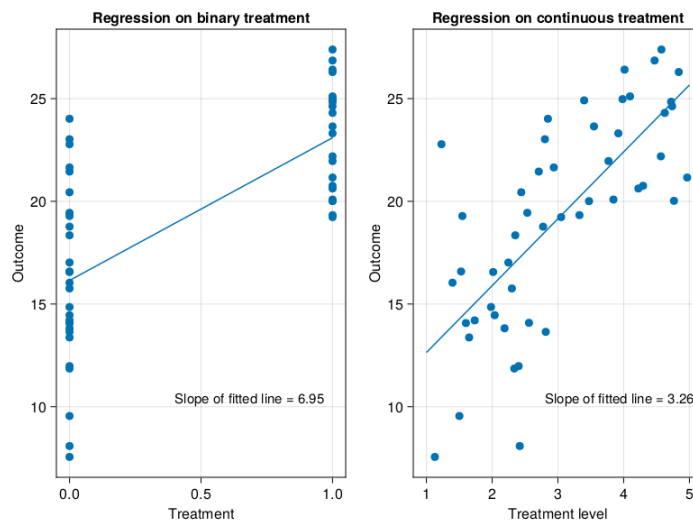
```

	parameters	mean	std	naive_se
1	:a	16.1438	0.6948	0.0109857
2	:b	6.95584	1.03417	0.0163517
3	: $\sigma$	3.66792	0.371859	0.00587961

```
• describe(chns1_2bt)
```

	parameters	median	mad_sd	mean	std
1	"a"	16.149	0.68	16.144	0.69
2	"b"	6.947	1.018	6.956	1.03
3	" $\sigma$ "	3.632	0.369	3.668	0.37

```
• begin
•   post1_2bt = DataFrame(chns1_2bt)[: , 3:5]
•   ms1_2bt = model_summary(post1_2bt,
•   names(post1_2bt))
• end
```



```

let
  x1 = 1.0:0.01:5.0
  f = Figure()
  medians = [ms1_2at[p, "median"] for p in
    [:a, :b, :σ]]
  ax = Axis(f[1, 2], title = "Regression
on continuous treatment",
    xlabel = "Treatment level", ylabel
    = "Outcome")
  sca1 = scatter!(x, y)
  annotations!("Slope of fitted line =
$(round(medians[2], digits=2))",
    position = (2.8, 10), fontsize=15)
  lin1 = lines!(x1, medians[1] .+
    medians[2] * x1)

  x2 = 0.0:0.01:1.0
  medians = [ms1_2bt[p, "median"] for p in
    [:a, :b, :σ]]
  ax = Axis(f[1, 1], title="Regression on
binary treatment",
    xlabel = "Treatment", ylabel =
    "Outcome")
  sca1 = scatter!(x_binary, y)
  lin1 = lines!(x2, medians[1] .+
    medians[2] * x2)
  annotations!("Slope of fitted line =
$(round(medians[2], digits=2))",
    position = (0.4, 10), fontsize=15)
f
end

```

ppl1\_3a (generic function with 2 methods)

```
• @model function ppl1_3a(x, y)
•   a ~ Normal(10, 5)
•   b ~ Normal(0, 5)
•   σ ~ Exponential(1)
•   μ = a .+ b .* x
•   for i in eachindex(x)
•     y[i] ~ Normal(μ[i], σ)
•   end
• end
```

ppl1\_3b (generic function with 2 methods)

```
• @model function ppl1_3b(x, y)
•   a ~ Normal(10, 5)
•   b_exp ~ Normal(5, 5)
•   σ ~ Exponential(1)
•   μ = a .+ b_exp .* exp.(-x)
•   for i in eachindex(x)
•     y[i] ~ Normal(μ[i], σ)
•   end
• end
```

```
• begin
•   #Random.seed!(1533)
•   n1 = 50
•   x1 = LinRange(1, 6, 50)
•   y1 = [rand(Normal(5 + 30exp(-x1[i]),
•   2), 1)[1] for i in 1:length(x1)]
• end;
```

	parameters	mean	std	naive_sd
1	:a	12.1371	0.96468	0.0152529
2	:b	-1.38376	0.255583	0.0040411
3	:σ	2.59658	0.262522	0.0041508

```
• begin
•   m1_3at = ppl1_3a(x1, y1)
•   chns1_3at = sample(m1_3at, NUTS(),
•   MCMCThreads(), 1000, 4)
•   describe(chns1_3at)
• end
```

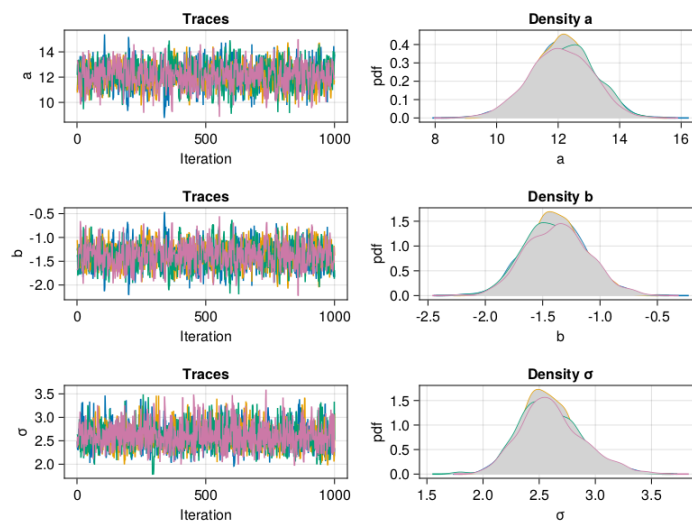


	parameters	median	mad_sd	mean	std
1	"a"	12.15	0.957	12.137	0.96
2	"b"	-1.384	0.257	-1.384	0.25
3	" $\sigma$ "	2.574	0.249	2.597	0.26

```

• begin
•   post1_3at = DataFrame(chns1_3at[:, :a,
•                         :b, : $\sigma$ ])
•   ms1_3at = model_summary(post1_3at, [:a,
•                                       :b, : $\sigma$ ])
• end


```



```

• plot_chains(post1_3at, [:a, :b, : $\sigma$ ])

```



```
• trankplot(post1_3at, "a")
```

	parameters	mean	std	naive_se
1	:a	5.80128	0.380655	0.00601869
2	:b_exp	20.26	2.92528	0.0462527
3	: $\sigma$	2.20733	0.234689	0.00371077

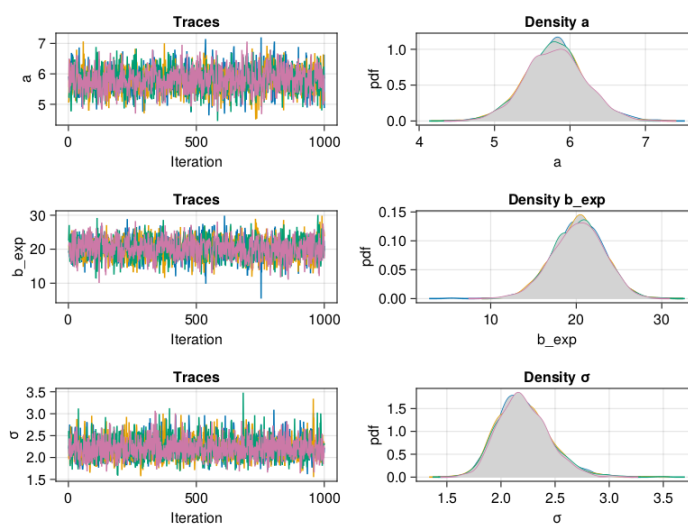
```
• begin
•   m1_3bt = ppl1_3b(x1, y1)
•   chns1_3bt = sample(m1_3bt, NUTS(),
•     MCMCThreads(), 1000, 4)
•   describe(chns1_3bt)
• end
```

	parameters	median	mad_sd	mean	std
1	"a"	5.802	0.368	5.801	0.38
2	"b_exp"	20.312	2.947	20.26	2.92
3	" $\sigma$ "	2.19	0.234	2.207	0.23

```

• begin
•   post1_3bt = DataFrame(chns1_3bt[:, :a,
•                         :b_exp, : $\sigma$ ])
•   ms1_3bt = model_summary(post1_3bt, [:a,
•                                       :b_exp, : $\sigma$ ])
• end


```



```

• plot_chains(post1_3bt, [:a, :b_exp, : $\sigma$ ])

```




- `trankplot(post1_3bt, "b_exp")`

► `[12.15, -1.384, 2.574]`

- `$\hat{a}_1$ ,  $\hat{b}$ ,  $\hat{\sigma}_1$  = ms1_3at[:, :median]`

► `[5.802, 20.312, 2.19]`

- `$\hat{a}_2$ ,  $\hat{b}_{exp}$ ,  $\hat{\sigma}_2$  = ms1_3bt[:, :median]`



```

• let
•   f = Figure()
•   ax = Axis(f[1, 1], title = "Linear
•       regression",
•       xlabel = "Treatments", ylabel =
•       "Outcomes")
•   scatter!(x1, y1)
•   lines!(x1,  $\hat{a}_1$  .+  $\hat{b}$  .* x1)
•
•   ax = Axis(f[2, 1], title = "Non-linear
•       regression",
•       xlabel = "Treatments", ylabel =
•       "Outcomes")
•   scatter!(x1, y1)
•   lines!(x1,  $\hat{a}_2$  .+  $\hat{b}_{exp}$  .* exp.(-x1))
•   f
• end

```

	xx	z	yy
1	3.1425	0	37.5294
2	0.0335661	1	30.0628
3	1.60408	0	28.1095
4	0.478735	1	31.3638
5	2.65874	0	35.7382
6	1.02705	1	36.0009
7	1.28799	0	24.3213
8	0.052966	1	29.5242
9	0.543994	0	25.4181
10	0.0304007	1	25.1863
: more			
100	0.00156342	1	28.8751

```

• begin
•   Random.seed!(12573)
•   n2 = 100
•   z = repeat([0, 1]; outer=50)
•   df1_8 = DataFrame()
•   df1_8.xx = [(z[i] == 0 ? rand(Normal(0,
•   1.2), 1).^2 : rand(Normal(0, 0.8),
•   1).^2)[1] for i in 1:n2]
•   df1_8.z = z
•   df1_8.yy = [rand(Normal(20 .+
•   5df1_8.xx[i] .+ 10df1_8.z[i], 3), 1)[1]
•   for i in 1:n2]
•   df1_8
• end

```

```
lm1_8 =
StatsModels.TableRegressionModel{LinearModel{GLM
```

```
yy ~ 1 + xx + z
```

Coefficients:

	Coef.	Std. Error	t	Pr(> t)
(Intercept)	20.1093	0.529823	37.95	<1e-16
xx	4.97503	0.213492	23.30	<1e-16
z	9.625	0.604978	15.91	<1e-16

```
• lm1_8 = lm(@formula(yy ~ xx + z), df1_8)
```

```
lm1_8_0 =
StatsModels.TableRegressionModel{LinearModel{GLM
```

```
yy ~ 1 + xx
```

Coefficients:

	Coef.	Std. Error	t	Pr(> t)
(Intercept)	20.0337	0.544062	36.82	<1e-16
xx	5.01957	0.226965	22.12	<1e-16

```
• lm1_8_0 = lm(@formula(yy ~ xx),
df1_8[df1_8.z .== 0, :])
```


```
lm1_8_1 =
StatsModels.TableRegressionModel{LinearModel{GLM
```

```
yy ~ 1 + xx
```

Coefficients:

	Coef.	Std. Error	t	Pr(> t)
(Intercept)	29.8841	0.49051	60.92	<1e-16
xx	4.66553	0.609796	7.65	<1e-16

```
• lm1_8_1 = lm(@formula(yy ~ xx),
df1_8[df1_8.z .== 1, :])
```



```

• let
•    $\hat{a}_1$ ,  $\hat{b}_1$  = coef(lm1_8_0)
•    $\hat{a}_2$ ,  $\hat{b}_2$  = coef(lm1_8_1)
•   x = range(0, maximum(df1_8.xx),
•             length=40)
•
•   f = Figure()
•   ax = Axis(f[1, 1]; title="Figure 1.8")
•   scatter!(df1_8.xx[df1_8.z .== 0],
•            df1_8.yy[df1_8.z .== 0])
•   scatter!(df1_8.xx[df1_8.z .== 1],
•            df1_8.yy[df1_8.z .== 1])
•   lines!(x,  $\hat{a}_1$  .+  $\hat{b}_1$  * x, label =
•           "Control")
•   lines!(x,  $\hat{a}_2$  .+  $\hat{b}_2$  * x, label =
•           "Treated")
•   axislegend(; position=(:right, :bottom))
•   current_figure()
• end

```

## 1.5 Classical and Bayesian inference.

No code.

## 1.6 Computing least-squares and Bayesian regression.

No code.



## 1.8 Exercises.

### Helicopters

```
helicopters =
```

	Helicopter_ID	width_cm	length_cm	time_s
<b>1</b>	1	4.6	8.2	1.64
<b>2</b>	1	4.6	8.2	1.74
<b>3</b>	1	4.6	8.2	1.68
<b>4</b>	1	4.6	8.2	1.62
<b>5</b>	1	4.6	8.2	1.68
<b>6</b>	1	4.6	8.2	1.7
<b>7</b>	1	4.6	8.2	1.62
<b>8</b>	1	4.6	8.2	1.66
<b>9</b>	1	4.6	8.2	1.69
<b>10</b>	1	4.6	8.2	1.62
⋮ more				
<b>20</b>	2	4.6	8.2	1.61

```
• helicopters =  
  CSV.read(ros_datadir("Helicopters",  
    "helicopters.csv"), DataFrame)
```

**Simulate 40 helicopters.**

	width_cm	length_cm	time_sec
1	4.32055	13.5698	1.92608
2	4.46568	12.2163	1.7315
3	2.23095	14.34	1.81414
4	6.69221	11.1536	1.69721
5	6.9636	11.7499	1.7386
6	0.464694	11.8161	1.57434
7	4.02125	9.37956	1.23015
8	1.84234	14.8711	1.65743
9	3.23692	7.12873	1.20158
10	2.92905	8.57974	1.36353
⋮	more		
40	5.69465	12.5771	1.88003

```

• begin
•   helis = DataFrame(width_cm =
•     rand(Normal(5, 2), 40), length_cm =
•     rand(Normal(10, 4), 40))
•   helis.time_sec = 0.5 .+ 0.04 .*
•   helis.width_cm .+ 0.08 .*
•   helis.length_cm .+ 0.1 .*
•   rand(Normal(0, 1), 40)
•   helis
• end

```

## Simulate 40 helicopters.

ppl1\_4 (generic function with 2 methods)

```

• @model function ppl1_4(w, l, y)
•   a ~ Normal(10, 5)
•   b ~ Normal(0, 5)
•   c ~ Normal(0, 5)
•   σ ~ Exponential(1)
•   μ = a .+ b .* w .+ c .* l
•   for i in eachindex(y)
•     y[i] ~ Normal(μ[i], σ)
•   end
• end

```

	parameters	mean	std	naiv
1	:a	0.523608	0.0660108	0.0010
2	:b	0.041239	0.00937394	0.0001
3	:c	0.0780014	0.00457018	7.2261
4	: $\sigma$	0.117155	0.0143296	0.0002

```

• begin
•   m1_4t = ppl1_4(helis.width_cm,
•   helis.length_cm, helis.time_sec)
•   chns1_4t = sample(m1_4t, NUTS(),
•   MCMCThreads(), 1000, 4)
•   describe(chns1_4t)
end

```

	parameters	median	mad_sd	mean	st
1	"a"	0.524	0.065	0.524	0.00
2	"b"	0.041	0.009	0.041	0.00
3	"c"	0.078	0.005	0.078	0.00
4	" $\sigma$ "	0.116	0.014	0.117	0.01

```

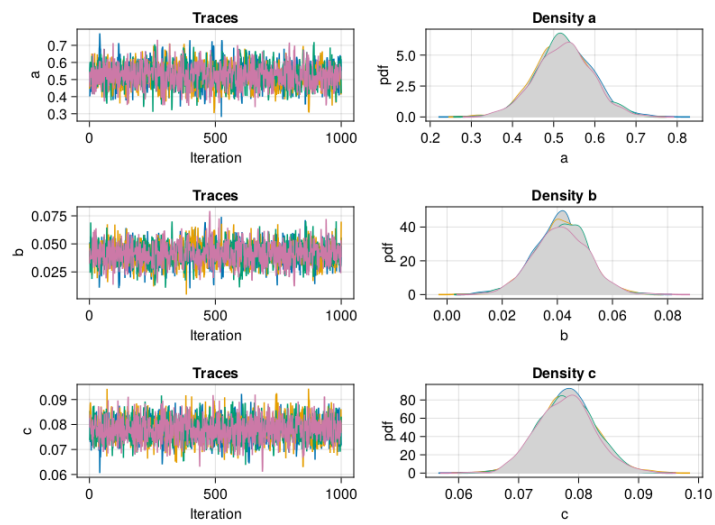
• begin
•   post1_4t = DataFrame(chns1_4t[:, :a, :b,
•   :c, : $\sigma$ ])
•   ms1_4t = model_summary(post1_4t, [:a,
•   :b, :c, : $\sigma$ ])
end

```


```

• ms1_4t[:, :b, :media]


```



```
• plot_chains(post1_4t, [:a, :b, :c])
```



```
• trunkplot(post1_4t, "b")
```



```

• let
•   w_range = LinRange(1.0, 8.0, 100)
•   w_times = mean.(link(post1_4t, (r, w) -
•   > r.a + r.c + r.b * w, w_range))
•
•   l_range = LinRange(6.0, 15.0, 100)
•   l_times = mean.(link(post1_4t, (r, l) -
•   > r.a + r.b + r.c * l, l_range))
•
•   f = Figure()
•   ax = Axis(f[1, 1], title = "Time in the
•   air on width and length",
•   xlabel = "Width/Length", ylabel =
•   "Time in the air")
•
•   lines!(w_range, w_times; label="Width")
•   lines!(l_range, l_times; label="Length")
•
•   f[1, 2] = Legend(f, ax, "Regression
•   lines", framevisible = false)
•
•   current_figure()
• end

```

## Note

Note that the `link` function is defined in both RegressionAndOtherStories (ROS) and Turing. In this case I added the `import` statement at the top of this notebook but I could also have qualified the call to `link ( ROS.link )`.

```
lnk1_4t =
```

```
▶ [[0.878075, 0.887118, 0.998595, 0.946481, 0.8760
```

```
• lnk1_4t = link(post1_4t, (r, l) -> r.a + r.b  
+ r.c * l, [5, 10, 12])
```

```
▶ [0.95517, 1.34411, 1.50066]
```

```
• median.(lnk1_4t)
```

```
▶ [0.0430268, 0.0372544, 0.038826]
```

```
• mad.(lnk1_4t)
```

```
▶ [0.954854, 1.34486, 1.50086]
```

```
• mean.(lnk1_4t)
```