

## See chapter 3 in Regression and Other Stories.

=====

### Widen the notebook

```
• html"""  
• <style>  
•     main {  
•         margin: 0 auto;  
•         max-width: 2000px;  
•         padding-left: max(160px, 10%);  
•         padding-right: max(160px, 10%);  
•     }  
• </style>  
• """
```

```
• using Pkg ✓ , DrWatson ✓
```

```
• begin  
•     # Specific to this notebook  
•     using GLM ✓  
•   
•     # Specific to ROSTuringPluto  
•     using Optim ✓  
•     using Logging ✓  
•     using Turing ✓  
•   
•     # Graphics related  
•     using GLMakie ✓  
•   
•     # Common data files and functions  
•     using RegressionAndOtherStories ✓  
•     import RegressionAndOtherStories: link  
•   
•     Logging.disable_logging(Logging.Warn)  
• end;
```

### 3.1 - Weighted averages

pop =

	stratum	country	population	avera
1	1	"United States"	310000000	36.8
2	2	"Mexico"	112000000	26.7
3	3	"Canada"	34000000	40.7

```
• pop = DataFrame(stratum=1:3, country=[  
  "United States", "Mexico", "Canada"],  
  population=Int[310e6, 112e6, 34e6],  
  average_age=[36.8, 26.7, 40.7])
```

34.61008771929824

```
• mean(pop.average_age,  
  weights(pop.population))
```


► [0.679825, 0.245614, 0.0745614]

```
• weights(pop.population)/sum(pop.population)
```

	variable	mean	min	median	
1	:stratum	2.0	1	2.0	3
2	:country	nothing	"Canada"	nothing	"
3	:population	1.52e8	34000000	1.12e8	3
4	:average_age	34.7333	26.7	36.8	4

```
• describe(pop)
```


## 3.3 - Graphing a line



```

• let
•   data = LinRange(0.01, 0.99, 200)
•   f = Figure(resolution = (800, 800))
•
•   for (i, scale) in enumerate([identity,
•   log10, log2, log, sqrt, Makie.logit])
•
•       row, col = fldmod1(i, 2)
•       Axis(f[row, col], yscale = scale,
•       title = string(scale),
•       yminorticksvisible = true,
•       yminorgridvisible = true,
•       yminorticks =
•       IntervalsBetween(8))
•
•       lines!(data, color = :blue)
•   end
•
•   f
• end


```



```

• let
•   data = 10 .^ LinRange(0.01, 5.0, 200)
•   f = Figure(resolution = (800, 300))
•
•   for (i, scale) in enumerate([log10,
•   log2, log])
•
•       row, col = fldmod1(i, 2)
•       Axis(f[1, i], yscale = scale, xscale
• = scale, title = string(scale),
•       yminorticksvisible = true,
•   yminorgridvisible = true,
•       yminorticks =
•   IntervalsBetween(8))
•
•       lines!(data, color = :blue)
•   end
•
•   f
• end

```



```

let
    f = Figure()
    ax = Axis(f[1, 1]; title = "Mile
record", subtitle = "(from 1900 to
2000)", xlabel = "Year", ylabel = "Time
[sec]")
    xlims!(ax, 1900, 2000)
    ax.xticks = 1900:25:2000
    x = 1900:2000
    y = 1007 .- 0.393 .* x
    lines!(x, y)
    annotations!("y = 1007 - 0.393 * x",
position=(1950, 250))
    current_figure()
end

```

## 3.4 - Log and exponential scales

Simulated data for metabolic.

	body_mass	rate
<b>1</b>	4.46877	4.72792
<b>2</b>	4.6314	4.77953
<b>3</b>	4.70212	4.86607
<b>4</b>	4.85071	5.09666
<b>5</b>	5.24315	5.23391
<b>6</b>	5.29702	5.27504
<b>7</b>	5.3163	5.42113
<b>8</b>	5.64758	5.54111
<b>9</b>	5.67783	5.54455
<b>10</b>	5.80551	5.76063
⋮ more		
<b>200</b>	9.20921	8.22882

```

• begin
•   x = sort(rand(Uniform(0.01, 10000),
•               200))
•   y = 4.1 * x.^0.74 .+ [rand.(Normal.(0,
•   sqrt(x[i])), 1)[1] for i in 1:length(x)]
•   metabolic = DataFrame(:body_mass => log.
•   (x), :rate => log.(y))
end

```

ppl3\_1 (generic function with 2 methods)

```

• @model function ppl3_1(m, r)
•   a ~ Normal(0, 0.3)
•   b ~ Normal(0, 0.3)
•   σ ~ Exponential(1)
•   μ = a .+ b .* m
•   for i in eachindex(y)
•     r[i] ~ Normal(μ[i], σ)
•   end
• end

```

	parameters	mean	std	naiv
1	:a	1.40073	0.0198207	0.0003
2	:b	0.740984	0.00240248	3.7986
3	: $\sigma$	0.0359071	0.00179637	2.8403

```

• begin
•   m3_1t = ppl3_1(metabolic.body_mass,
•                 metabolic.rate)
•   chns3_1t = sample(m3_1t, NUTS(),
•                     MCMCThreads(), 1000, 4)
•   describe(chns3_1t)
end


```

	parameters	median	mad_sd	mean	st
1	"a"	1.401	0.02	1.401	0.02
2	"b"	0.741	0.002	0.741	0.00
3	" $\sigma$ "	0.036	0.002	0.036	0.00

```

• begin
•   post3_1t = DataFrame(chns3_1t)[: , 3:5]
•   ms3_1t = model_summary(post3_1t,
•                           names(post3_1t))
end

```



```

• let
•   x = LinRange(1, 10000, 1000)
•   y = 4.1 * x.^0.74
•   f = Figure()
•   ax = Axis(f[1, 1]; title="Metabolic
•   rate (linear scale)", xlabel="Body mass
•   [kg]", ylabel="Metobolic rate [W]")
•   scatter!(exp.(metabolic.body_mass), exp.
•   (metabolic.rate))
•   lines!(x, y; color=:darkred)

•   ax = Axis(f[1, 2]; title="Metabolic
•   rate (log-log scale)", xscale=log,
•   yscale=log,
•   xlabel="log(body mass [kg])",
•   ylabel="log(metobolic rate [W])")
•   x =
•   LinRange(minimum(metabolic.body_mass),
•   maximum(metabolic.body_mass), 100)
•   scatter!(metabolic.body_mass,
•   metabolic.rate)
•   lines!(x, ms3_1t["a", "median"] .+
•   ms3_1t["b", "median"] * x;
•   color=:darkred)
•   current_figure()

end

```

```

100-element LinRange{Float64, Int64}:
4.46877, 4.51665, 4.56454, 4.61242, 4.6603, ..., 9

```

```

• LinRange(minimum(metabolic.body_mass),
• maximum(metabolic.body_mass), 100)

```

```
57.69669926958961
```

```
• exp(exp(1.4))
```




## 3.5 - Probability distributions

	height	sex
1	61.1316	"female"
2	64.6008	"female"
3	67.182	"female"
4	64.6839	"female"
5	61.9961	"female"
6	61.2885	"female"
7	58.3058	"female"
8	64.6222	"female"
9	66.6702	"female"
10	65.9953	"female"
⋮ more		
200000	73.614	"male"

```
• begin
•   N = 100000
•   heights = DataFrame()
•   height = vcat(rand(Normal(63.7, 2.7),
•   N),
•   rand(Normal(69.1, 2.9), N))
•   sex = repeat(["female", "male"],
•   inner=N)
•   heights.height = height
•   heights.sex = sex
•   heights
• end
```

```
► (mean = 63.703, var = 7.28943, std = 2.69989, m
```

```
• begin
•   menHeights = heights[heights.sex ==
•   "male", :height]
•   womenHeights = heights[heights.sex ==
•   "female", :height]
•   (mean=mean(womenHeights),
•   var=var(womenHeights),
•   std=std(womenHeights),
•   median=median(womenHeights),
•   mad_sd=mad(womenHeights))
end
```



```
• let
•   f = Figure()
•   ax = Axis(f[1, 1]; title="Density
•   heights")
•   density!(womenHeights; color=color
•   = (:lightgreen, 0.4), label="Women")
•   density!(menHeights; color=color =
•   (:lightblue, 0.4), label="Men")
•   axislegend()
•   f
end
```

```
0.49714220980937984
```


```
• begin
•   wdf = Normal(63.65, 2.68)
•   cdf(wdf, 63.65 + 0.67 * 2.68) -
•   cdf(wdf, 63.65 - 0.67 * 2.68)
end
```

0.6826894921370859

- $\text{cdf}(\text{wdf}, 63.65 + 2.68) - \text{cdf}(\text{wdf}, 63.65 - 2.68)$

0.9544997361036417

- $\text{cdf}(\text{wdf}, 63.65 + 2 * 2.68) - \text{cdf}(\text{wdf}, 63.65 - 2 * 2.68)$



```

• let
•   wdf = Normal(63.65, 2.68)
•   x = range(55.0, 72.5 ; length=100)
•   lines(x, pdf.(wdf, x); color=:darkblue)
•
•   x1 = range(63.65 - 3 * 2.68, 63.65 - 2
•             * 2.68; length=20)
•   band!(x1, fill(0, length(x1)), pdf.
•         (wdf, x1);
•         color = (:blue, 0.25), label =
•             "Label")
•
•   x1 = range(63.65 + 2 * 2.68, 63.65 + 3
•             * 2.68; length=20)
•   band!(x1, fill(0, length(x1)), pdf.
•         (wdf, x1);
•         color = (:blue, 0.25), label =
•             "Label")
•
•   x1 = range(63.65 - 2 * 2.68, 63.65 - 1
•             * 2.68; length=20)
•   band!(x1, fill(0, length(x1)), pdf.
•         (wdf, x1);
•         color = (:blue, 0.45), label =
•             "Label")
•
•   x1 = range(63.65 + 1 * 2.68, 63.65 + 2
•             * 2.68; length=20)
•   band!(x1, fill(0, length(x1)), pdf.
•         (wdf, x1);
•         color = (:blue, 0.45), label =
•             "Label")
•
•   x1 = range(63.65 - 1 * 2.68, 63.65;
•             length=20)
•
•

```

```

• band!(x1, fill(0, length(x1)), pdf.
• (wdf, x1);
• color = (:blue, 0.55), label =
• "Label")
•

x1 = range(63.65, 63.65 + 2.68;
length=20)
band!(x1, fill(0, length(x1)), pdf.
(wdf, x1);
color = (:blue, 0.55), label =
"Label")

text!("68%", position = (63.65, 0.05),
align = (:center, :center),
textsize = 30)
text!("13.5%", position = (67.5, 0.02),
align = (:center, :center),
textsize = 20)
text!("13.5%", position = (59.6, 0.02),
align = (:center, :center),
textsize = 20)
text!("2.5%", position = (69.75,
0.0045), align = (:center, :center),
textsize = 15)
text!("2.5%", position = (57.7,
0.0045), align = (:center, :center),
textsize = 15)
current_figure()

end

```

► ( $\hat{m} = 6.0014$ ,  $m = 6.0$ ,  $\hat{\sigma} = 2.04294$ ,  $\sigma = 2.04939$ )

```

• let
• n = 20; p = 0.3
• y = rand(Binomial(n, p), 10000)
• ( $\hat{m} = \text{mean}(y)$ ,  $m = 20 * 0.3$ ,  $\hat{\sigma} = \text{std}(y)$ ,
•  $\sigma = \text{sqrt}(n * p * (1 - p))$ )
end

```


► ( $\hat{m} = 0.295$ ,  $m = 0.3$ )

```

• let
• n = 20; p = 0.3
• y = rand(Bernoulli(p), 10000)
• ( $\hat{m} = \text{mean}(y)$ ,  $m = 0.3$ )
• end

```

## LogNormal



```

• let
•   menw = rand(LogNormal(5.13, 0.17),
•   10000)
•   menwl = log.(menw)
•   f = Figure()
•   ax = Axis(f[1, 1]; title="Log weights
•   of men\n(Normal distribution)",
•   xlabel="Logarithm of weight [lb]",
•   ylabel="Density")
•   density!(menwl)
•   ax = Axis(f[1, 2]; title="Weights of
•   men\n(LogNormal distribution",
•   xlabel="Weight of men [lb]",
•   ylabel="Density")
•   density!(menw)
•   current_figure()
end

```

## Binomial

	parameters	median	mad_sd	mean	st
1	"bv"	6.0	1.483	6.0	2.07

```

• let
•   df = DataFrame(bv = rand(Binomial(20,
•   0.3), 1000))
•   model_summary(df, [:bv])
end

```

```
► (mean = 6.0, std = 2.04939)
```

```
• begin  
•   n = 20  
•   p = 0.3  
•   (mean = n * p, std =  $\sqrt{n * p * (1 - p)}$ )  
• end
```

## Poisson

```
► [5, 5, 3, 4, 4, 9, 2, 3, 2, 8]
```

```
• rand(Poisson(4.52), 10)
```

## 3.6 - Probability modeling

```
20689.577579211084
```

```
• 1 / (pdf(Normal(0.49, 0.04), 0.5) / 200000)
```

```
20.689577579211083
```

```
• 1 / (1000pdf(Normal(0.49, 0.04), 0.5) /  
200000)
```