

## See Chapter 6 in Regression and Other Stories.

Widen the notebook.

```
• html"""  
• <style>  
•     main {  
•         margin: 0 auto;  
•         max-width: 2000px;  
•         padding-left: max(160px, 10%);  
•         padding-right: max(160px, 10%);  
•     }  
• </style>  
• """
```

```
• using Pkg ✓ , DrWatson ✓
```

A typical set of Julia packages to include in notebooks.

```
• begin  
•     # Specific to this notebook  
•     using GLM ✓  
•   
•     # Specific to ROSTuringPluto  
•     using Optim ✓  
•     using Logging ✓  
•     using Turing ✓  
•   
•     # Graphics related  
•     using GLMakie ✓  
•   
•     # Common data files and functions  
•     using RegressionAndOtherStories ✓  
•     import RegressionAndOtherStories: link  
•   
•     Logging.disable_logging(Logging.Warn)  
• end;
```

## 6.1 Regression models.

## 6.2 Fitting a simple regression to fake data.

	x	y
1	1.0	0.175189
2	2.0	1.70123
3	3.0	1.22936
4	4.0	0.79866
5	5.0	1.31689
6	6.0	2.30556
7	7.0	3.02729
8	8.0	3.24706
9	9.0	2.04172
10	10.0	2.74121
⋮	more	
20	20.0	6.72016

```
• let
•   n = 20
•   x = LinRange(1, n, 20)
•   a = 0.2
•   b = 0.3
•   sigma = 0.5
•   y = a .+ b .* x .+ rand(Normal(0,
•   sigma), n)
•   global fake = DataFrame(x=x, y=y)
end
```

ppl6\_1 (generic function with 2 methods)


```
• @model function ppl6_1(x, y)
•   a ~ Uniform(-2, 2)
•   b ~ Uniform(-2, 2)
•   σ ~ Uniform(0, 10)
•   μ = a .+ b .* x
•   for i in eachindex(y)
•     y[i] ~ Normal(μ[i], σ)
•   end
• end
```

► [	parameters	mean	std	naive_
1	:a	0.203286	0.281582	0.004452
2	:b	0.302298	0.0232234	0.000367
3	:σ	0.603167	0.110236	0.001742

```
• begin
•   m6_1t = ppl6_1(fake.x, fake.y)
•   chns6_1t = sample(m6_1t, NUTS(),
•   MCMCThreads(), 1000, 4)
•   describe(chns6_1t)
• end
```

	parameters	median	mad_sd	mean	st
1	"a"	0.199	0.279	0.203	0.28
2	"b"	0.303	0.023	0.302	0.02
3	"σ"	0.587	0.102	0.603	0.11

```
• begin
•   post6_1t = DataFrame(chns6_1t)[: , 3:5]
•   ms6_1t = model_summary(post6_1t,
•   names(post6_1t))
• end
```



```

let
  f = Figure()

  ax = Axis(f[1, 1]; title="Regression of
fake data.", xlabel="fake.x",
ylabel="fake.y")
  scatter!(fake.x, fake.y)
  x = 1:0.01:20
  y = ms6_1t["a", "median"] .+
ms6_1t["b", "median"] .* x
  lines!(x, y)
  â, b̂, ε̂ = round.(ms6_1t[:, "median"];
digits=2)
  annotations!("y = $(â) + $(b̂) * x + ε";
position=(5, 0.8))

  ax = Axis(f[1, 2]; title="Regression of
fake data.", subtitle="(using the
link() function)",
xlabel="fake.x", ylabel="fake.y")
  scatter!(fake.x, fake.y)
  xrange = LinRange(1, 20, 200)
  y = mean.(link(post6_1t, (r,x) -> r.a +
x * r.b, xrange))
  lines!(xrange, y)
  annotations!("y = $(â) + $(b̂) * x + ε";
position=(5, 0.8))

  current_figure()
end

```

	parameters	simulated	median	mad_sd
1	:a	0.2	0.199	0.279
2	:b	0.3	0.303	0.023
3	: $\sigma$	0.5	0.587	0.102

```
• DataFrame(parameters = Symbol.
  (names(post6_1t)), simulated = [0.2, 0.3,
  0.5], median = ms6_1t[:, "median"], mad_sd
  = ms6_1t[:, "mad_sd"])
```

## 6.3 Interpret coefficients as comparisons, not effects.

	earnk	height	male
1	50.0	74	1
2	60.0	66	0
3	30.0	64	0
4	25.0	65	0
5	50.0	63	0
6	62.0	68	0
7	51.0	63	0
8	9.0	64	0
9	29.0	62	0
10	32.0	73	1
⋮ more			
1816	6.0	68	1

```

• begin
•   earnings =
•   CSV.read(ros_datadir("Earnings",
•   "earnings.csv"), DataFrame)
•   earnings[:, [:earnk, :height, :male]]
end

```

	variable	mean	min	median	max	n
1	:earnk	21.1473	0.0	16.0	400.0	0
2	:height	66.5688	57	66.0	82	0
3	:male	0.371696	0	0.0	1	0

```

• describe(earnings[:, [:earnk, :height,
• :male]])

```

ppl6\_2 (generic function with 2 methods)

```
• @model function ppl6_2(male, height, earnk)
•   a ~ Normal()
•   b ~ Normal()
•   c ~ Normal()
•   σ ~ Exponential(1)
•   μ = a .+ b .* height .+ c .* male
•   for i in eachindex(earnk)
•     earnk[i] ~ Normal(μ[i], σ)
•   end
• end
```

```
► [ parameters      mean      std      naive

1  :a      -0.676695  0.995374  0.01573
2  :b       0.297586  0.0173254 0.00027
3  :c       5.89894   0.741121  0.01171
4  :σ      21.5295    0.354245  0.00566
```

```
• begin
•   m6_2t = ppl6_2(earnings.male,
•   earnings.height, earnings.earnk)
•   chns6_2t = sample(m6_2t, NUTS(),
•   MCMCThreads(), 1000, 4)
•   describe(chns6_2t)
• end
```

```
parameters  median  mad_sd  mean  st

1  "a"      -0.689   1.01    -0.677  0.99
2  "b"       0.298   0.017   0.298   0.01
3  "c"       5.911   0.736   5.899   0.74
4  "σ"      21.524   0.348   21.53   0.35
```

```
• begin
•   post6_2t = DataFrame(chns6_2t)[: , 3:6]
•   ms6_2t = model_summary(post6_2t,
•   names(post6_2t))
• end
```



```

let
  â, b̂, ĉ, ô = round.(ms6_2t[:,
    "median"]; digits=2)

  fig = Figure()

  ax = Axis(fig[1, 1]; title="Earnings
    for males", subtitle="earnk = $(round(ĉ
    + â; digits=2)) + $(b̂) * mheight + ε")
  m = sort(earnings[earnings.male .== 1,
    [:height, :earnk]])
  scatter!(m.height, m.earnk)
  mheight_range =
    LinRange(minimum(m.height),
    maximum(m.height), 200)
  earnk = mean.(link(post6_2t, (r,x) ->
    r.c + r.a + x * r.b, mheight_range))

  lines!(mheight_range, earnk;
    color=:darkred)

  ax = Axis(fig[1, 2]; title="Earnings
    for females", subtitle="earnk = $(â) +
    $(b̂) * fheight + ε")
  f = sort(earnings[earnings.male .== 0,
    [:height, :earnk]])
  scatter!(f.height, f.earnk)
  fheight_range =
    LinRange(minimum(f.height),
    maximum(f.height), 200)
  earnk = mean.(link(post6_2t, (r,x) ->
    r.a + x * r.b, fheight_range))
  lines!(fheight_range, earnk;
    color=:darkred)

  fig

```



```
end
```

```
R2 = 0.08694354978286623
```

```
• R2 = 1 - ms6_2t["σ", "mean"]^2 /  
  std(earnings.earnk)^2
```

## 6.4 Historical origins of regression.

```
heights =
```

	daughter_height	mother_height
1	52.5	59.5
2	52.5	59.5
3	53.5	59.5
4	53.5	59.5
5	55.5	59.5
6	55.5	59.5
7	55.5	59.5
8	55.5	59.5
9	56.5	58.5
10	56.5	58.5
⋮ more		
5524	73.5	63.5

```
• heights =  
  CSV.read(ros_datadir("PearsonLee",  
    "heights.csv"), DataFrame)
```

ppl6\_3 (generic function with 2 methods)


```
• @model function ppl6_3(m_height, d_height)
•   a ~ Normal()
•   b ~ Normal()
•   σ ~ Exponential(1)
•   μ = a .+ b .* m_height
•   for i in eachindex(d_height)
•     d_height[i] ~ Normal(μ[i], σ)
•   end
• end
```

► [	parameters	mean	std	naive_
1	:a	18.0761	0.627493	0.009921
2	:b	0.732217	0.0100345	0.000158
3	:σ	2.30724	0.0223726	0.000353

```
• begin
•   m6_3t = ppl6_3(heights.mother_height,
•   heights.daughter_height)
•   chns6_3t = sample(m6_3t, NUTS(),
•   MCMCThreads(), 1000, 4)
•   describe(chns6_3t)
• end
```

	parameters	median	mad_sd	mean	st
1	"a"	18.087	0.614	18.076	0.62
2	"b"	0.732	0.01	0.732	0.01
3	"σ"	2.307	0.022	2.307	0.02

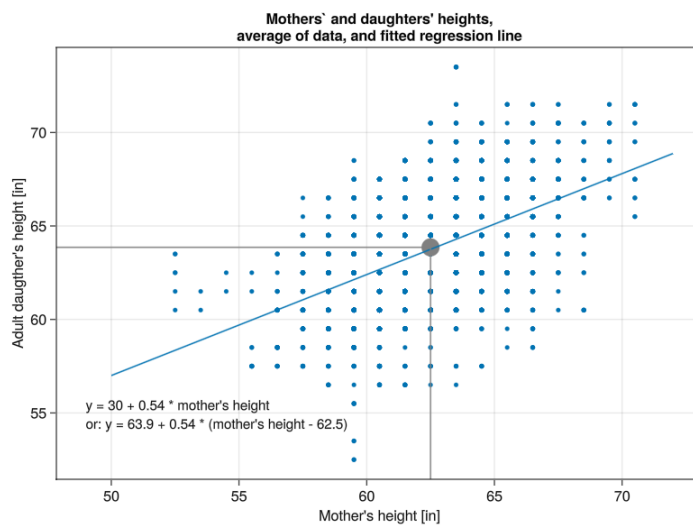
```
• begin
•   post6_3t = DataFrame(chns6_3t)[: , 3:5]
•   ms6_3t = model_summary(post6_3t,
•   names(post6_3t))
• end
```



```

• let
•   f = Figure()
•   ax = Axis(f[1, 1]; title="Mothers' and
•   daughters' heights")
•   xlims!(ax, 51, 74)
•   scatter!(jitter.
•   (heights.mother_height), jitter.
•   (heights.daughter_height); markersize=3)
•   x_range = LinRange(51, 74, 100)
•   lines!(x_range, mean.(link(post6_3t,
•   (r, x) -> r.a + r.b * x, x_range));
•   color=:darkred)
•   scatter!([mean(heights.mother_height)],
•   [mean(heights.daughter_height)];
•   markersize=20)
•   f
end


```



```

let
    f = Figure()
    ax = Axis(f[1, 1]; title="Mothers' and
    daughters' heights,\naverage of data,
    and fitted regression line",
        xlabel="Mother's height [in]",
        ylabel="Adult daughter's height
        [in]")
    scatter!(heights.mother_height,
    heights.daughter_height; markersize=5)
    xrange = LinRange(50, 72, 100)
    y = 30 .+ 0.54 .* xrange
    m̄ = mean(heights.mother_height)
    d̄ = mean(heights.daughter_height)
    scatter!([m̄], [d̄]; markersize=20,
    color=:gray)
    lines!(xrange, y)
    vlines!(ax, m̄; ymax=0.55, color=:grey)
    hlines!(ax, d̄; xmax=0.58, color=:grey)
    annotations!("y = 30 + 0.54 * mother's
    height", position=(49, 55), fontsize=15)
    annotations!("or: y = 63.9 + 0.54 *
    (mother's height - 62.5)", position=
    (49, 54), fontsize=15)
    f
end

```



```

• let
•   f = Figure()
•   ax = Axis(f[1, 1]; title="Mothers' and
•       daughters' heights,\naverage of data,
•       and fitted regression line",
•       xlabel="Mother's height [in]",
•       ylabel="Adult daughter's height
•           [in]")
•   scatter!(heights.mother_height,
•             heights.daughter_height; markersize=5)
•   xrange = LinRange(0, 72, 100)
•   y = 30 .+ 0.54 .* xrange
•   m̄ = mean(heights.mother_height)
•   d̄ = mean(heights.daughter_height)
•   scatter!([m̄], [d̄]; markersize=20,
•           color=:gray)
•   lines!(xrange, y)
•   annotations!("y = 30 + 0.54 * mother's
• height", position=(20, 35), fontsize=15)
•   annotations!("or: y = 63.9 + 0.54 *
• (mother's height - 62.5)", position=
• (20, 33), fontsize=15)
•   f
• end

```

ppl6\_4 (generic function with 2 methods)

```

• @model function ppl6_4(m_height, d_height)
•   a ~ Normal(25, 3)
•   b ~ Normal(0, 0.5)
•   σ ~ Exponential(1)
•   μ = a .+ b .* m_height
•   for i in eachindex(d_height)
•       d_height[i] ~ Normal(μ[i], σ)
•   end
• end

```

	parameters	mean	std	naive_s
1	:a	29.4845	0.760218	0.0120201
2	:b	0.54995	0.0121652	0.0001923
3	: $\sigma$	2.26237	0.0216619	0.0003425

```

• begin
•   m6_4t = ppl6_4(heights.mother_height,
•                 heights.daughter_height)
•   chns6_4t = sample(m6_4t, NUTS(),
•                     MCMCThreads(), 1000, 4)
•   describe(chns6_4t)
end

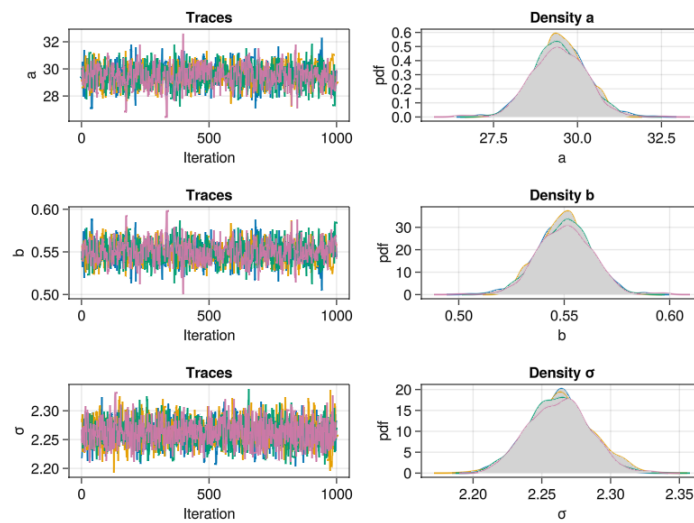
```

	parameters	median	mad_sd	mean	std
1	"a"	29.473	0.734	29.485	0.760
2	"b"	0.55	0.012	0.55	0.012
3	" $\sigma$ "	2.262	0.022	2.262	0.022


```

• begin
•   post6_4t = DataFrame(chns6_4t)[: , 3:5]
•   ms6_4t = model_summary(post6_4t,
•                          names(post6_4t))
end

```



```
• plot_chains(post6_4t, [:a, :b, :sigma])
```



```
• trunkplot(post6_4t, "b")
```

Above trunkplot and the low ess numbers a couple of cells earlier do not look healthy.

## 6.5 The paradox of regression to the mean.

	midterm	final
<b>1</b>	70.8711	55.178
<b>2</b>	21.7327	44.2992
<b>3</b>	67.0099	53.225
<b>4</b>	27.4136	50.4825
<b>5</b>	42.3918	68.7687
<b>6</b>	80.811	81.8612
<b>7</b>	34.2082	52.7463
<b>8</b>	52.8503	47.5438
<b>9</b>	62.3295	69.732
<b>10</b>	63.4512	69.7803
: more		
<b>1000</b>	72.4838	65.602

```

• let
•   n = 1000
•   true_ability = rand(Normal(50, 10), n)
•   noise_1 = rand(Normal(0, 10), n)
•   noise_2 = rand(Normal(0, 10), n)
•   midterm = true_ability + noise_1
•   final = true_ability + noise_2
•   global exams =
•   DataFrame(midterm=midterm, final=final)
end

```

ppl6\_5 (generic function with 2 methods)

```

• @model function ppl6_5(midterm, final)
•   a ~ Normal()
•   b ~ Normal()
•   σ ~ Exponential(1)
•   μ = a .+ b .* midterm
•   for i in eachindex(final)
•       final[i] ~ Normal(μ[i], σ)
•   end
• end

```



	parameters	mean	std	naive_
1	:a	7.31704	0.870082	0.013757
2	:b	0.830372	0.0178971	0.000282
3	: $\sigma$	13.4318	0.307599	0.004863

```

• begin
•   m6_5t = ppl6_5(exams.midterm,
•     exams.final)
•   chns6_5t = sample(m6_5t, NUTS(),
•     MCMCThreads(), 1000, 4)
•   describe(chns6_5t)
end

```

	parameters	median	mad_sd	mean	st
1	"a"	7.31	0.866	7.317	0.87
2	"b"	0.831	0.018	0.83	0.01
3	" $\sigma$ "	13.43	0.308	13.432	0.30

```


• begin
•   post6_5t = DataFrame(chns6_5t)[: , 3:5]
•   ms6_5t = model_summary(post6_5t,
•     names(post6_5t))
end

```

```
df_poll =
```

	poll1	poll2	poll3	poll4	poll5
<b>1</b>	2002	10.0	70.0	25.0	5.0
<b>2</b>	2002	5.0	72.0	25.0	3.0
<b>3</b>	2001	10.0	68.0	26.0	6.0
<b>4</b>	2001	5.0	65.0	27.0	8.0
<b>5</b>	2001	2.0	67.0	25.0	8.0
<b>6</b>	2000	8.0	67.0	28.0	5.0
<b>7</b>	2000	6.0	66.0	26.0	8.0
<b>8</b>	2000	2.0	66.0	28.0	6.0
<b>9</b>	1999	5.0	71.0	22.0	7.0
<b>10</b>	1995	9.0	77.0	13.0	10.0
⋮ more					
<b>32</b>	1937	12.0	60.0	33.0	7.0

```
• df_poll = CSV.read(ros_datadir("Death",  
  "polls.csv"), DataFrame)
```



```

let
  f = Figure()
  ax = Axis(f[1, 1]; title="Death penalty
  opinions", xlabel="Year",
  ylabel="Percentage support for the
  death penalty")
  scatter!(df_poll.year, df_poll.support
  .* 100)
  err_lims = [100(sqrt(df_poll.support[i]*
  (1-df_poll.support[i])/1000)) for i in
  1:nrow(df_poll)]
  errorbars!(df_poll.year, df_poll.support
  .* 100, err_lims, color = :red)
  f
end

```

Used in later notebooks.

	STATE	TOTLDF	DOR	DORAVG	HRS
1	"AL"	296.0	33.47	32.65	11.61
2	"AR"	77.0	15.4	15.65	9.7
3	"AZ"	231.0	41.5	39.42	7.92
4	"CA"	528.0	9.21	9.14	8.8
5	"FL"	851.0	30.19	30.18	10.91
6	"GA"	323.0	19.63	19.12	12.78
7	"ID"	31.0	48.48	44.16	3.55
8	"IL"	238.0	11.26	10.98	8.18
9	"IN"	79.0	11.81	10.93	5.61
10	"KY"	59.0	10.67	10.24	7.03
	⋮ more				
26	"WY"	5.0	9.98	11.63	4.58

```

• begin
•   death_raw=CSV.read(ros_datadir("Death",
•     "dataforandy.csv"), DataFrame;
•     missingstring="NA")
•   death =
•     death_raw[completecases(death_raw), :]
end

```

```

• let
•   st_abbr = death[:, 1]
•   ex_rate = death[:, 8] ./ 100
•   err_rate = death[:, 7] ./ 100
•   hom_rate = death[:, 5] ./ 100000
•   ds_per_homicide = death[:, 3] ./ 1000
•   ds = death[:, 2]
•   hom = ds ./ ds_per_homicide
•   ex = ex_rate .* ds
•   err = err_rate .* ds
•   pop = hom ./ hom_rate
•   std_err_rate = sqrt.( (err .+ 1) .* (ds
•     .+ 1 .- err) ./ ((ds .+ 2).^2 .* (ds .+
•     3)) )
• end;

```

